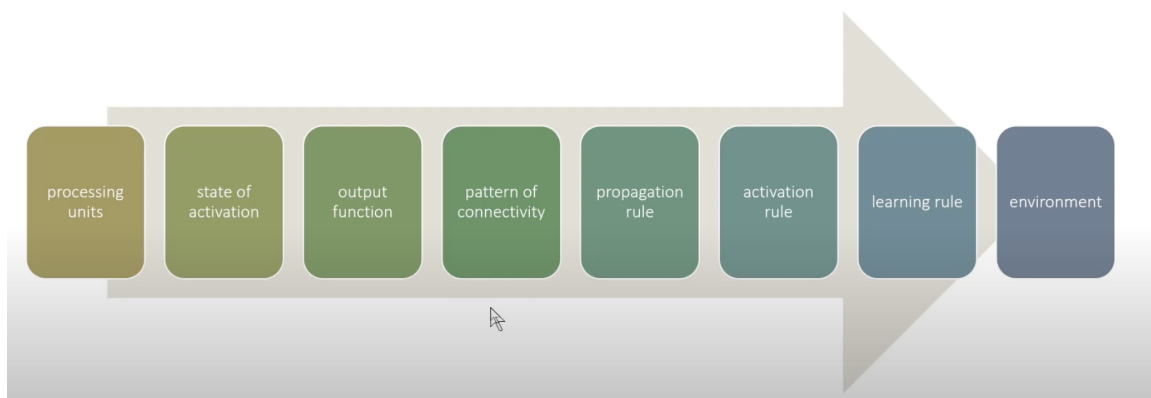


Lesson 1

- Started in 1943, with a mathematical model of a neuron
- 1953: Mark 1 Perceptron
- AI Winter: Minsky and Papert wrote a book that said a single layer of perceptrons couldn't model an XOR gate. Multiple layers could, but this was largely not noticed.
- 1986: Parallel Distributed Processing




- Adding just one extra layer of neurons is enough to allow any mathematical model to be approximated with a neural network.
 - In practice, slow and big
 - to get practical performance, they require even more layers of neurons
 - Deep instead of wide


-
- Software stack: fastai → PyTorch → Python
 - fastai is designed for teaching + development
 - layered API
 - Ability to switch to Keras + TF
 - Remember to discuss on: forums.fastai.com
 - Randomness involved in training a model

- Don't expect same results every time

- In regular programming




- doesn't work very well for recognising pictures
- Arthur Samuel pioneered: instead of telling the exact steps required to solve a problem, show examples of the problem to solve and let it figure it out itself




Suppose we arrange for some automatic means of testing the effectiveness of any current weight assignment in terms of actual performance and provide a mechanism for altering the weight assignment so as to maximize the performance. We need not go into the details of such a procedure to see that it could be made entirely automatic and to see that a machine so programmed would “learn” from its experience.

- Not just depend on inputs, but also a set of weights (parameters to alter the way to solve the problem), then we can alter the weights to find out the best way to solve the solution



- Requires a function in *model* that can approximate any problem using different set of weights depending on the weights. Neural network is that function. *Universal Approximation Theorem* proves that a neural network can be used to solve any problem to any level of accuracy.
- A general way is required to update the weights. This is Stochastic Gradient Descent (SGD).
- The terminology has been updated



- functional form of model → architecture
 - weights → parameters
 - predictions are calculated from independent variables (data without labels)
 - results → predictions
 - measure of performance → loss
-
- Learning approach only creates predictions, not recommended actions
 - data needs labels
 - Positive feedback loop
 - original data is biased making the model biased
 - more the model is used, more biased the data becomes

- making the model even more biased
- Error rate is always calculated on validation dataset
 - to prevent overfitting
 - if error rate were calculated on training dataset, it would be optimised for that dataset and not work well on other dataset