

# Chapter 6

## 1. How could multi-label classification improve the usability of the bear classifier?

Allow for the prediction of no bears present in the image. Otherwise, a multi-class classifier will always predict a class of bear that it thinks is most likely.

## 2. How do we encode the dependent variable in a multi-label classification problem?

As a one-hot encoded vector.

## 3. How do you access the rows and columns of a DataFrame as if it was a matrix?

To access the m-th row and the n-th column: `df.iloc[m,n]`

## 4. How do you get a column by name from a DataFrame

```
df['column_name']
```

## 5. What is the difference between a dataset and DataLoader?

`Dataset`: tuple of independent and dependent variable for a *single* item

`DataLoader`: iterator that provides a stream of mini-batches, where each mini-batch is a tuple of a batch of independent variables and dependent variables (but doesn't have to be; you can pass any collection to a `DataLoader` to be split into mini-batches)

## 6. What does a Datasets object normally contain?

A training `Dataset` and a validation `Dataset`

## 7. What does a DataLoaders object normally contain?

A training `DataLoader` and a validation `DataLoader`

## 8. What does lambda do in Python?

`Lambda` are shortcuts for writing one-line functions.

## 9. What are the methods to customise how the independent and dependent variables are created with the data block API?

`get_x` is used to specify how the independent variables are created

`get_y` is used to specify how the dependent variables are created

## 10. Why is softmax not an appropriate output activation function when using a one hot encoded target?

One-hot encoded vectors are used when the target label has more than one positive labels, aka multi-label classifications. However, `softmax` makes it such that the model will select one class over the others, making it ideal for training a classifier when each item has one definite label, but not an appropriate activation function for multi-label classification.

## 11. Why is nll\_loss not an appropriate loss function when using a one hot encoded target?

`nll_loss` can only output and select one class over the others. Hence, it can't be used for multi-label classification where one-hot encoded targets are used.

## 12. What is the difference between nn.BCELoss and nn.BCEWithLogitsLoss?

`nn.BCELoss` does not include the initial sigmoid. It assumes that the appropriate activation function (ie. the sigmoid) has already been applied to the predictions. `nn.BCEWithLogitsLoss`, on the other hand, does both the sigmoid and cross entropy in a single function.

## 13. Why can't we use regular accuracy in a multi-label problem?

Regular accuracy when the final class predicted (and only one of it), is the one with the highest activation. Multi-label however requires comparing the

sigmoid of the activations with a threshold as multiple classes can be applicable to an item.

## 14. When is it okay to tune an hyper-parameter on the validation set?

It is okay to do so when the relationship between the hyper-parameter and the metric being observed is smooth. With such a smooth relationship, we would not be picking an inappropriate outlier.

## 15. How is `y_range` implemented in fastai?

`y_range` is implemented using `sigmoid_range` in fastai:

```
def sigmoid_range(x, low, high):  
    return x.sigmoid() * (high - low) + low
```

## 16. What is a regression problem? What loss function should you use for such a problem?

A problem where the dependent variable are continuous values. They use `nn.MSELoss` or Mean Squared Error loss function.

## 17. What do you need to do to make sure the fastai library applies the same data augmentation to your input images and your target point coordinates?

Use `blocks = (ImageBlock, PointBlock)` in the `DataBlock`. This makes sure that fastai automatically and correctly applies data augmentation to coordinates.